Demo document — not a scientific publication

ClipFormer: a transformer baseline at two operating points

E. Fixture, F. Sample — Department of Invented Vision

Abstract
ClipFormer is a plain transformer baseline for action recognition. We train on the DemoKinetics-200 training split with supervised labels and report two operating points rather than a single configuration.

3. Results

Table 2 reports both configurations on the ED-Bench validation split. ClipFormer-B with a ViT-B backbone and 8 frames per clip reaches 77.0% top-1 accuracy. ClipFormer-L with a ViT-L backbone and 16 frames per clip reaches 83.1% top-1 accuracy.

We do not report inference latency in this version of the paper.

4. Limitations

Both configurations were trained with the same recipe, so the comparison between them is internally consistent but says nothing about runtime cost.